

# SYCL II

Soner Steiner

Intel certified oneAPI Instructor  
sonersteiner at gmail dot com

# Overview

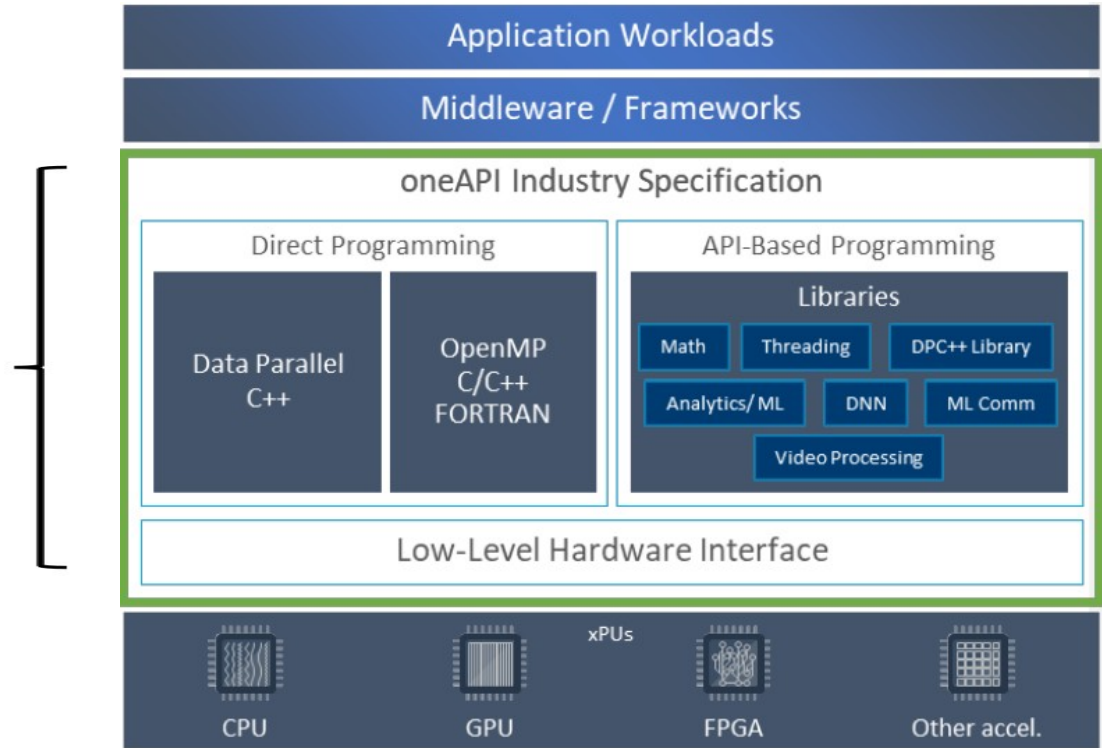
- Manage the data transfer  
Buffers and Unified Shared Memory
- Basic parallel kernels
- ND-Range kernels
- Sub-groups
- Reductions

# Programmers' perspective: Three things to consider

1. Offload the code to  
device

2. Manage the transfer  
of Data

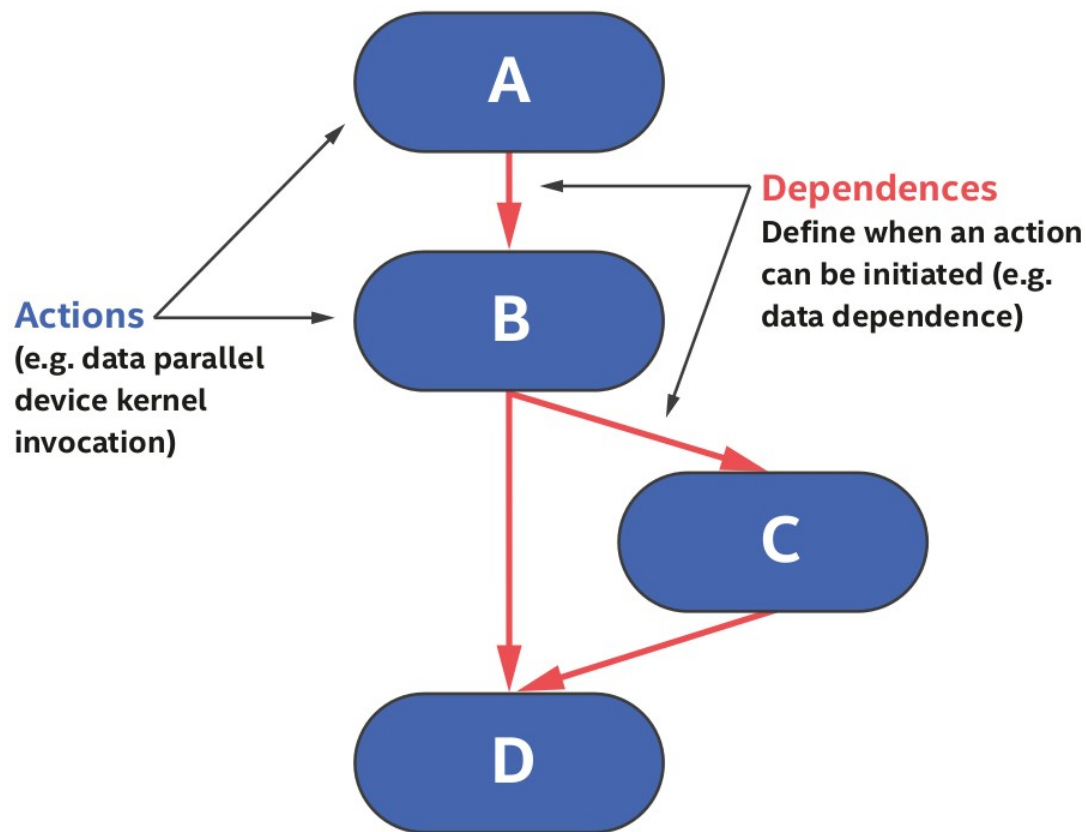
3. Implement  
Parallelism



# Memory Models

- **Buffer Memory Model** – abstract view of memory that can be local to the host or a device, and is accessible via accessors.
- **Unified Shared Memory (USM)**- pointer-based approach for memory model that is familiar for C++ programmers.
- **Images**: a special type of buffer that has an extra functionality specific to image processing

# Task Graphs (Directed Acyclic Graph)



- **Dependency** resolution and node **execution** are controlled by the **runtime**
- **Dependencies** determine the order that **kernels** are **executed** in
- Dependencies can be **explicit** or **implicit**

# Explicit Dependencies Using Events

```
constexpr int N = 101;
int main()
{
    queue q;
    int *data = malloc_shared<int>(N, q);

    auto e = q.parallel_for(N, [=] (id<1> i) { data[i] = i ;} );
    q.submit( [&] (handler &h)
    {
        h.depends_on(e);
        h.single_task([=] ()
        {
            for(int i = 1; i < N; ++i)
                data[0] += data[i];
        } );
    } );
    q.wait();

    std::cout << "printing sum after computation \n" ;
    std::cout << data[0] << " ";
    std::cout << "\n" ;
}
```

- Create event to initialize the data in kernel1
- Kernel2 sums up the elements
- 5050

# Buffer Memory Model

**Buffers** encapsulate data shared between host and device

**Accessors** provide access to data stored in buffers and create data dependencies in the graph.

**Unified Shared Memory (USM)** provides an alternative pointer-based mechanism for managing memory

```
queue q;  
std::vector<int> v(N, 3);  
{  
    buffer buf(v);  
    q.submit( [&] (handler& h)  
    {  
        accessor a(buf, h, write_only);  
        h.parallel_for(N, [=] (auto i) { a[i] = i; } );  
    } );  
}  
  
for (int i = 0; i < N; i++) std::cout << v[i] << " ";
```

## Buffer Creation – two approaches

- Construct a new buffer using `sycl::range` to specify the size, data will not be initialized!

```
buffer( const sycl::range<dimensions> &bufferRange,  
        const sycl::property_list &proplist={} );
```

- Create buffer from existing data, data will be copied!

```
Buffer( T, hostData,  
        const sycl::range<dimensions> &bufferRange,  
        const sycl::property_list &proplist={} );
```



# Examples of Buffer Creation

```
buffer b1{v};  
buffer b2{v.begin(), v.end()};
```

Buffer for vectors

```
// create a buffer of ints from std::array  
std::array<int, 42> data;  
buffer b3{data};
```

Buffer for std::array

```
// create a buffer of 5 doubles and initialize it from  
// a host pointer  
double dd[5] = {1.1, 2.2, 3.14, 4.4, 5.5};  
buffer b4{dd, range{5}};
```

Buffer from a host pointer

```
std::cout << "printing v before computation \n" ;  
for (int i = 0; i < N; i++) std::cout << v[i] << " ";  
std::cout << "\n" ;
```

# Accessors

- Only means of **accessing data** in Buffers!
- They create the **dependencies** for the **runtime**.

# Accessor Modes

Access Mode	Description
<code>read_only</code>	Read only Access
<code>write_only</code>	Write-only accessor Previous Contents not discarded
<code>read_write</code>	Read and Write access

# Code Walkthrough

```
#include <CL/sycl.hpp>
using namespace sycl;

int main() {
    std::vector<float> A(1024, 1.0f), B(1024, 2.0f), C(1024);
    {
        buffer bufA {A}, bufB {B}, bufC {C};
        queue q;
        q.submit([&](handler &h) {
            auto A = bufA.get_access(h, read_only);
            auto B = bufB.get_access(h, read_only);
            auto C = bufC.get_access(h, write_only);
            h.parallel_for(1024, [=](auto i) {
                C[i] = A[i] + B[i];
            });
        });
    }
}
```

# Host Accessor

(up to now our accessors have been in the command group)

- The **Host Accessor** is an accessor which uses host buffer access target.
- Host accessors make **data available for access on the host**.
- They **synchronize with the host** by defining a new dependence between the currently accessing graph and the host.
- Creating host accessor is a **blocking call**.

# Some Dependency Patterns

# Linear Dependence Using In-order queue

Create In-order  
queue

Initialize the data in  
Kernel 1

Kernel 2 sums up  
the elements

```
constexpr int N=42;

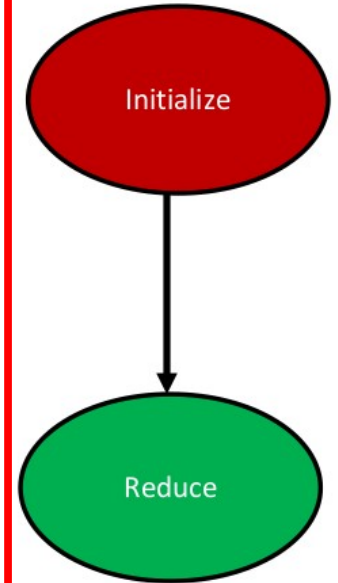
int main()
{
    queue Q{property::queue::in_order()};

    int *data = malloc_shared<int>(N,Q);

    Q.parallel_for(N, [=](id<1> i) { data[i] = 1; });

    Q.single_task([=]()
    {
        for(int i=1; i < N; ++i)
            data[0] += data[i];
    });
    Q.wait();

    assert(data[0] == N);
    for(int i = 0; i < N; ++i)
        std::cout << data[i] << " ";
    std::cout << "\n";
    return 0;
}
```



# Linear Dependence Using Buffers and Accessors

Use Buffers and  
Accessors to Initialize  
the data in Kernel1

```
constexpr int N=101;
int main()
{
    queue q;
    buffer <int> data{ range{N} };

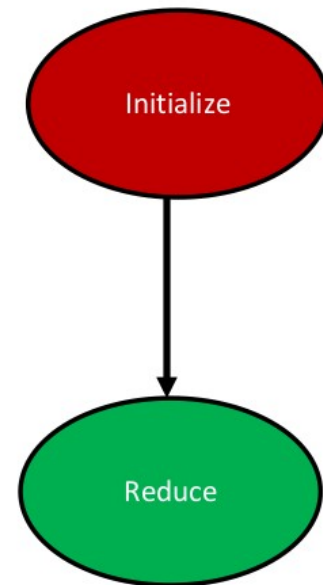
    q.submit( [&] (handler &h)
    {
        accessor a{data, h};
        h.parallel_for(N, [=] (id<1> i) { a[i] = i; } );
    } );

    q.submit( [&] (handler &h)
    {
        accessor a{data, h};
        h.single_task([=] ()
        {
            for(int i = 1; i < N; ++i)
                a[0] += a[i];
        } );
    } );

    host_accessor h_a{data};
    std::cout << h_a[0] << "\n";

    return 0;
}
```

Kernel 2 sums up  
the elements



5050



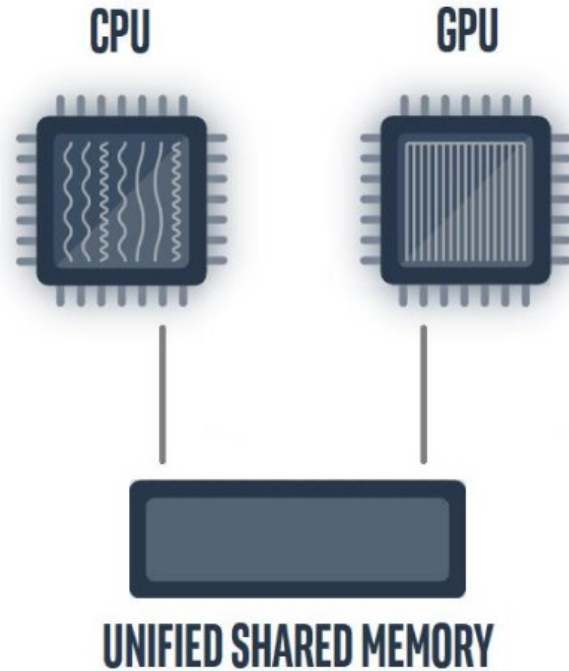
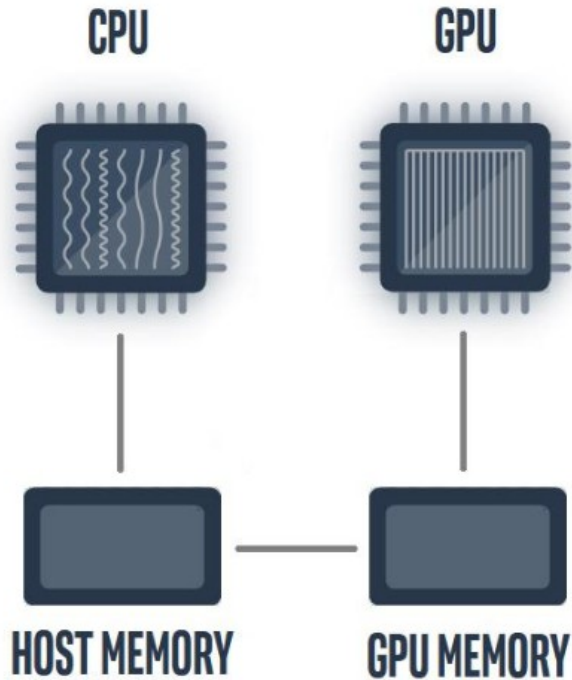
# Unified shared memory (USM)

USM provides a pointer-based alternative in SYCL

- Simplifies porting to an accelerator
- Gives programmers the desired level of control
- Complementary to buffers

# Developer View of USM

Developers can reference the **same memory object** in host and device code with USM



# Unified shared memory (USM)

USM provides both **explicit** and **implicit** models for managing memory.

Allocation Type	Description	Accessible on HOST	Accessible on DEVICE
device	Allocations in device memory ( <b>explicit</b> )	NO	YES
host	Allocations in host memory ( <b>implicit</b> )	YES	YES
shared	Allocations can migrate between host and device memory ( <b>implicit</b> )	YES	YES

*Automatic data accessibility and explicit data movement supported.*

# USM – Explicit Data Movement

```
queue q;  
int hostArray[N];  
int *deviceArray = (int*) malloc_device(N * sizeof(int), q);  
  
for(int i = 0; i < N; ++i) hostArray[i] = i;  
  
// copy hostArray to deviceArray  
q.memcpy(deviceArray, &hostArray[0], N*sizeof(int));  
q.wait();  
  
q.submit( [&] (handler &h)  
{  
    h.parallel_for(N, [=] (auto ID)  
    {  
        deviceArray[ID] = ID*ID ;  
    } );  
} );  
q.wait();  
  
//copy deviceArray back to hostArray  
q.memcpy(&hostArray[0], deviceArray, N*sizeof(int));  
q.wait();  
free(deviceArray, q);
```

malloc\_device

mem\_copy

mem\_copy

# USM – Implicit Data Movement

```
queue q;  
int *hostArray = (int*) malloc_host(N * sizeof(int), q);  
int *sharedArray = (int*) malloc_shared(N * sizeof(int), q);  
  
for(int i = 0; i < N; ++i) hostArray[i] = i;  
  
q.submit( [&] (handler &h)  
{  
    h.parallel_for(N, [=] (auto ID)  
    {  
        sharedArray[ID] = hostArray[ID] * hostArray[ID];  
    } );  
} );  
q.wait();  
  
for (int i = 0; i < N; i++) hostArray[i] = sharedArray[i] ;  
free(hostArray, q);  
free(sharedArray, q);
```

malloc\_host  
malloc\_shared

# Unified Shared Memory – When to use it?

## SYCL\* **Buffers are powerful** and elegant

- Use if the abstraction applies cleanly in your application, and/or buffers aren't disruptive to your development

## USM provides a familiar pointer-based C++ interface

- Useful when **porting C++ code** to SYCL, by minimizing changes
- Use shared allocations when porting code, **to get functional quickly**
- Note that shared allocation is **not intended** to provide peak performance out of box
- Use explicit USM allocations when **controlled data movement** is needed

# USM – Data Dependency in Queues

No accessors in USM

Dependencies must be specified explicitly using events

- `queue.wait()`
- wait on `event` objects
- use the `depends_on` method inside a command group

# USM – Data Dependency in Queues

```
queue q;  
int *data = (int*) malloc_shared(N * sizeof(int), q);  
for(int i = 0; i < N; ++i) data[i] = i;  
  
q.submit( [&] (handler &h)  
{  
    h.parallel_for<class taskA>(range<1> (N), [=] (id<1> i)  
    {  
        data[i] += 1;  
    } );  
} );  
q.wait();  
q.submit( [&] (handler &h)  
{  
    h.parallel_for<class taskB>(range<1> (N), [=] (id<1> i)  
    {  
        data[i] += 2;  
    } );  
} );  
q.wait();  
q.submit( [&] (handler &h)  
{  
    h.parallel_for<class taskC>(range<1> (N), [=] (id<1> i)  
    {  
        data[i] += 3;  
    } );  
} );  
q.wait();  
  
for (int i = 0; i < N; i++) std::cout << data[i] << " ";  
free(data, q);
```

Explicit **wait()** used to ensure  
Data dependency is  
maintained

**wait()** will **block execution** on  
host





# USM – Data Dependency in Queues

```
queue q;
int* data = malloc_shared<int>(N, q);

for(int i = 0; i < N; ++i) data[i] = i;

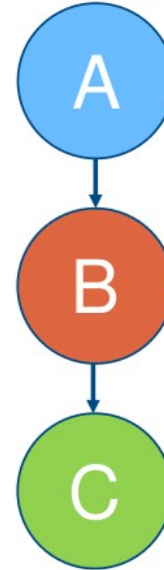
auto e1 = q.submit( [&] (handler &h)
{
    h.parallel_for<class taskA>(range<1>(N), [=] (id<1> i)
    {
        data[i] += 1;
    });
});

auto e2 = q.submit( [&] (handler &h)
{
    h.depends_on(e1);
    h.parallel_for<class taskB>(range<1>(N), [=] (id<1> i)
    {
        data[i] += 2;
    });
});

// non-blocking: execution of host code is possible
q.submit( [&] (handler &h)
{
    h.depends_on(e2);
    h.parallel_for<class taskC>(range<1>(N), [=] (id<1> i)
    {
        data[i] += 3;
    });
});

q.wait();
std::cout << "printing data after computation \n" ;
for (int i = 0; i < N; i++) std::cout << data[i] << " ";
free(data, q);
```

use **depends\_on** method to let command group handler know that specified event should be complete before specified task can execute.



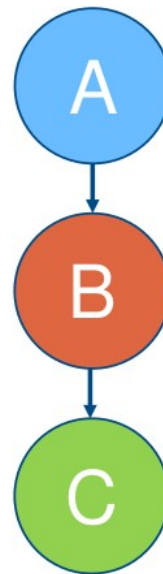
# USM – Data Dependency in Queues

```
queue q{property::queue::in_order()};
int* data = malloc_shared<int>(N, q);

for(int i = 0; i < N; ++i) data[i] = i;
q.submit( [&] (handler &h)
{
    h.parallel_for<class taskA>(range<1> (N), [=] (id<1> i)
    {
        data[i] += 1;
    } );
} );
q.submit( [&] (handler &h)
{
    h.parallel_for<class taskB>(range<1> (N), [=] (id<1> i)
    {
        data[i] += 2;
    } );
} );
q.submit( [&] (handler &h)
{
    h.parallel_for<class taskC>(range<1> (N), [=] (id<1> i)
    {
        data[i] += 3;
    } );
} );
q.wait();
free(data, q);
```

use **in\_queue** property for the queue

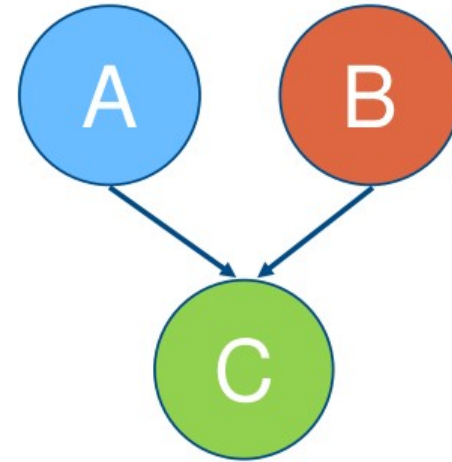
Execution will not overlap even  
If the queues have no data dependency



# USM – Data Dependency in Queues

```
queue q;  
int *data1 = (int*) malloc_shared(N * sizeof(int), q);  
int *data2 = (int*) malloc_shared(N * sizeof(int), q);  
  
for(int i = 0; i < N; ++i){ data1[i] = 10; data2[i] = 20;}  
  
auto e1 = q.submit( [&] (handler &h)  
{  
    h.parallel_for<class taskA>(range<1> (N), [=] (id<1> i)  
    {  
        data1[i] += 1;  
    } );  
});  
  
auto e2 = q.submit( [&] (handler &h)  
{  
    h.parallel_for<class taskB>(range<1> (N), [=] (id<1> i)  
    {  
        data2[i] += 2;  
    } );  
});  
  
q.submit( [&] (handler &h)  
{  
    h.depends_on({e1, e2});  
    h.parallel_for<class taskC>(range<1> (N), [=] (id<1> i)  
    {  
        data1[i] += data2[i];  
    } );  
});  
  
q.wait();  
for (int i = 0; i < N; i++) std::cout << data1[i] << " ";  
free(data1, q); free(data2, q);
```

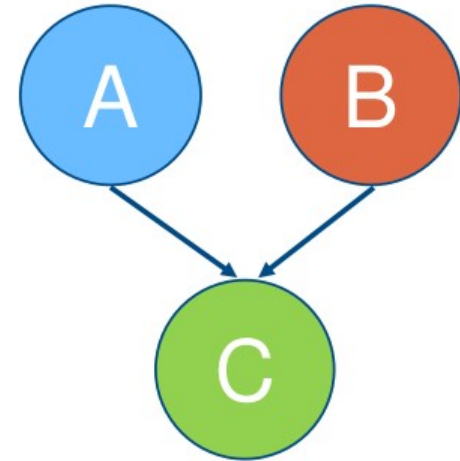
use `depends_on()` method to let command group handler know that specified events should be complete before specified tasks can execute.



# USM – Data Dependency in Queues

```
queue q;  
int* data1 = malloc_shared<int>(N, q);  
int* data2 = malloc_shared<int>(N, q);  
  
for(int i = 0; i < N; ++i){ data1[i] = 10; data2[i] = 20;}  
  
auto e1 = q.parallel_for<class taskA>(range<1>(N), [=] (id<1> i)  
{  
    data1[i] += 1;  
} );  
auto e2 = q.parallel_for<class taskB>(range<1>(N), [=] (id<1> i)  
{  
    data2[i] += 2;  
} );  
q.parallel_for<class taskC>(range<1>(N), {e1, e2}, [=] (id<1> i)  
{  
    data1[i] += data2[i];  
} );  
q.wait();  
free(data1, q); free(data2, q);
```

A more **simplified** way of specifying dependency as parameter of `parallel_for`



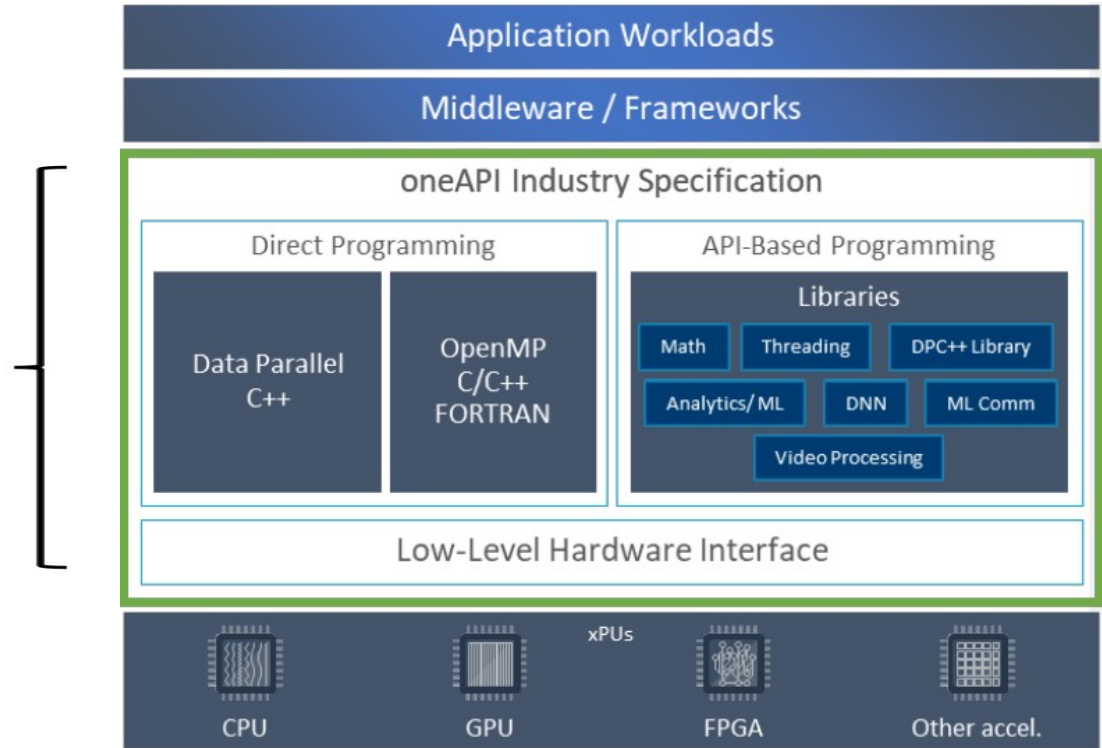
# Unified Shared Memory

- Summary

- What is Unified Shared Memory (USM)?
- Implicit and Explicit data movement between host and device
- Handling data dependency in multiple kernel tasks using wait event, depends\_on method and in\_order queue property

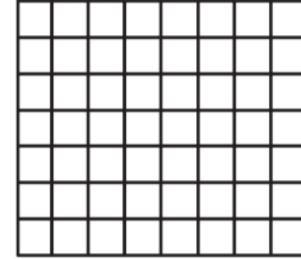
# Programmers' perspective: Three things to consider

1. Offload the code to device
2. Manage the transfer of Data
3. Implement Parallelism

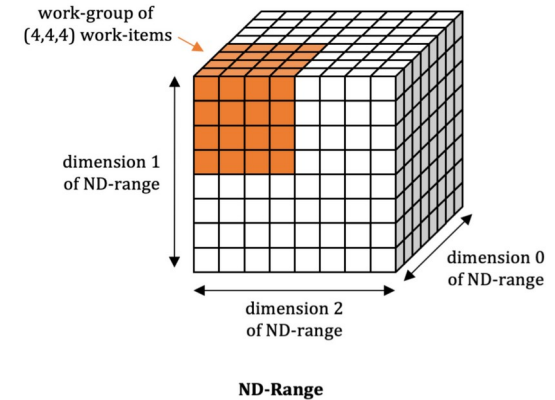


# Three forms of Parallel Kernels

- **Basic Parallel Kernels**



- **ND-range Parallel Kernels**



- **Hierarchical Parallel Kernels ('Experimental alternative syntax')**

# Basic Parallel Kernels

- Parallel kernel allows multiple instances of an operation to execute in parallel.
- Useful to offload parallel execution of a basic `for-loop` in which each iteration is completely independent and in any order.
- Parallel kernels are expressed using the `parallel_for` function.
- Up to the programmer to handle/confirm that there are no `dependencies`.

for-loop in CPU application

```
for(int i = 0; i < N; ++i)
{
    c[i] = a[i] + c[i] ;
}
```

Offload to a accelerator using `parallel_for`

```
h.parallel_for(range<1>(N), [=](id<1> i)
{
    C[i] = A[i] + B[i] ;
});
```



# Basic Parallel Kernels

The functionality of basic parallel kernels is exposed via `range`, `id` and `item` classes

- `range` class is used to describe the iteration space of parallel execution
- `id` class is used to index an individual instance of a kernel in a parallel execution

```
h.parallel_for(range<1>(N), [=](id<1> idx)
{
    //CODE THAT RUNS ON DEVICE
});
```

# Basic Parallel Kernels

The functionality of basic parallel kernels is exposed via **range**, **id** and **item** classes

- **range** class is used to describe the iteration space of parallel execution
- **id** class is used to index an individual instance of a kernel in a parallel execution
- **item** class represents an individual instance of a kernel function, exposes additional functions to query properties of the execution range

```
h.parallel_for(range<1>(N), [=](id<1> idx)
{
    //CODE THAT RUNS ON DEVICE
});
```

```
h.parallel_for(range<1>(N), [=](item<1> item)
{
    auto idx = item.get_id();
    auto R    = item.get_range();
    |
    //CODE THAT RUNS ON DEVICE
});
```

# Basic Parallel Kernels

The functionality of basic parallel kernels is exposed via `range`, `id` and `item` classes

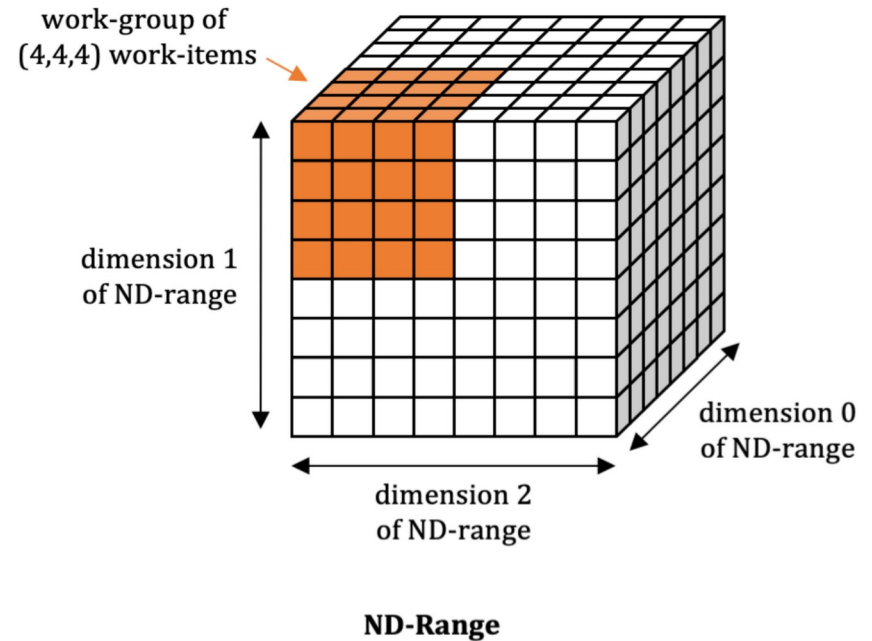
- **Dimensionality**  
<1>, <2> or <3>  
is templated and must be declared at COMPILE time
- **Size** is dynamic – passed to constructor at runtime

```
h.parallel_for(range<1>(N), [=](id<1> idx)
{
    //CODE THAT RUNS ON DEVICE
});
```

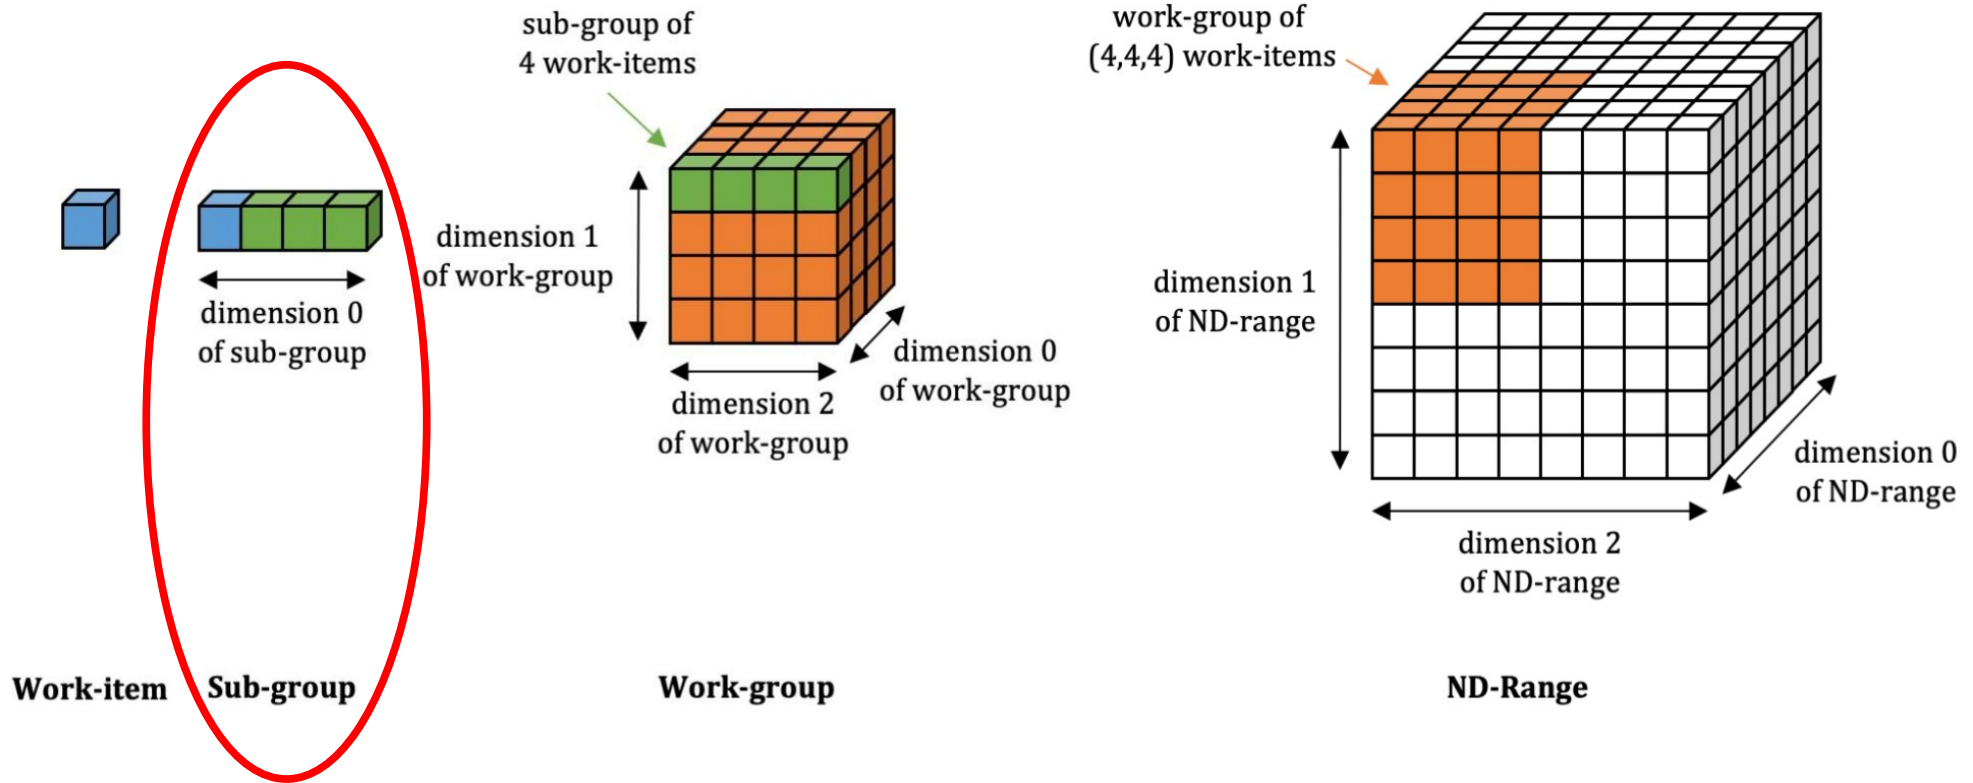
```
h.parallel_for(range<1>(N), [=](item<1> item)
{
    auto idx = item.get_id();
    auto R    = item.get_range();
    |
    //CODE THAT RUNS ON DEVICE
});
```

# ND-range Kernels

- **ND-range kernels** enable **low level performance tuning** by providing access to local memory and mapping executions to compute units on hardware.
- The entire iteration space is divided into smaller groups called **work-groups**, work-items within a work-group are scheduled on a single compute unit on hardware.
- The grouping of kernel executions into work-groups will allow control of **resource usage** and **load balance** work distribution.

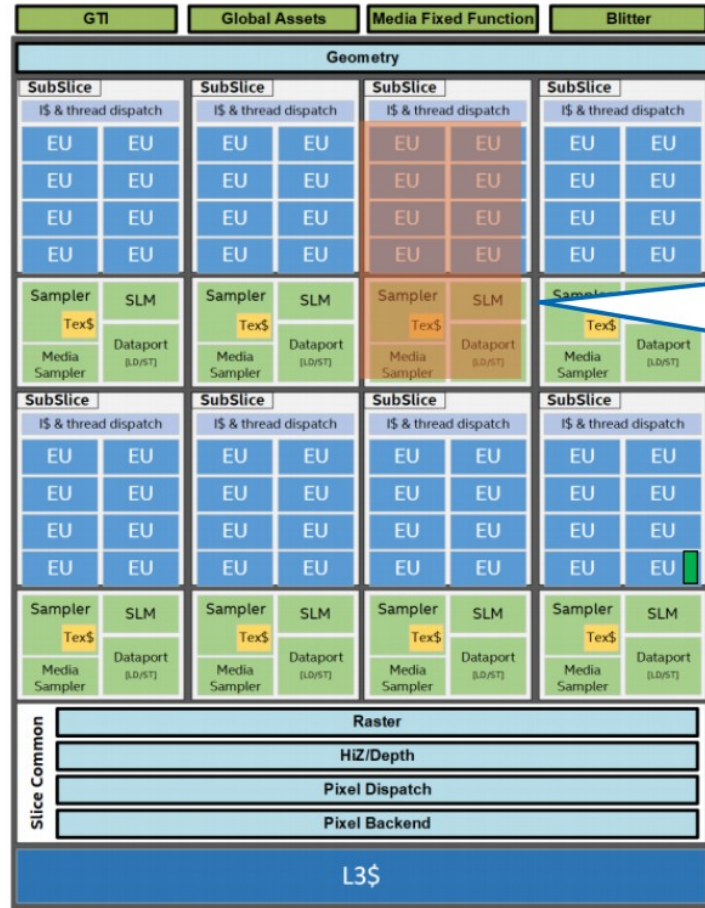


# SYCL Thread Hierarchy and Mapping

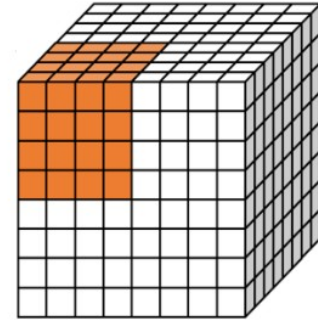


Covered later

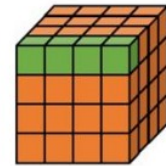
# SYLC Thread Hierarchy and Mapping



All work-items in a **work-group** are scheduled on one Compute Unit, which has its own local memory

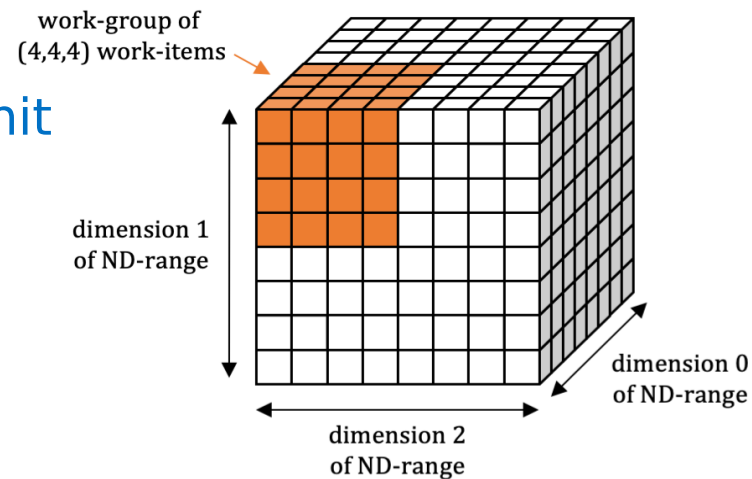


All work-items in a **sub-group** are mapped to vector hardware



# ND-range Kernels

- Basic Parallel Kernels are easy way to parallelize a for-loop but **does not allow** performance optimization at hardware level.
- **ND-range kernel** is another way to express parallelism which enable **low level performance tuning** by providing access to local memory and mapping executions compute units on hardware.
- The entire iteration space is divided into smaller groups called **work-groups**, work-items within a work-group are scheduled on a single compute unit on hardware.
- The grouping of kernel executions into work-groups will allow control of **resource usage** and **load balance** work distribution.



# ND-range Kernels

The functionality of `nd_range` kernels is exposed via `nd_range` and `nd_item` classes

`nd_range` class represents a grouped execution range using global execution range and the local execution range of each work-group.

`nd_item` class represents an individual instance of a kernel function and allows to query for work-group range and index.



# Sub-groups

# Sub-groups

Understand how **Sub-Groups** map to GPU hardware

Understand how using **Sub-Groups shuffle operations** can achieve better performance and avoid repeated global memory access

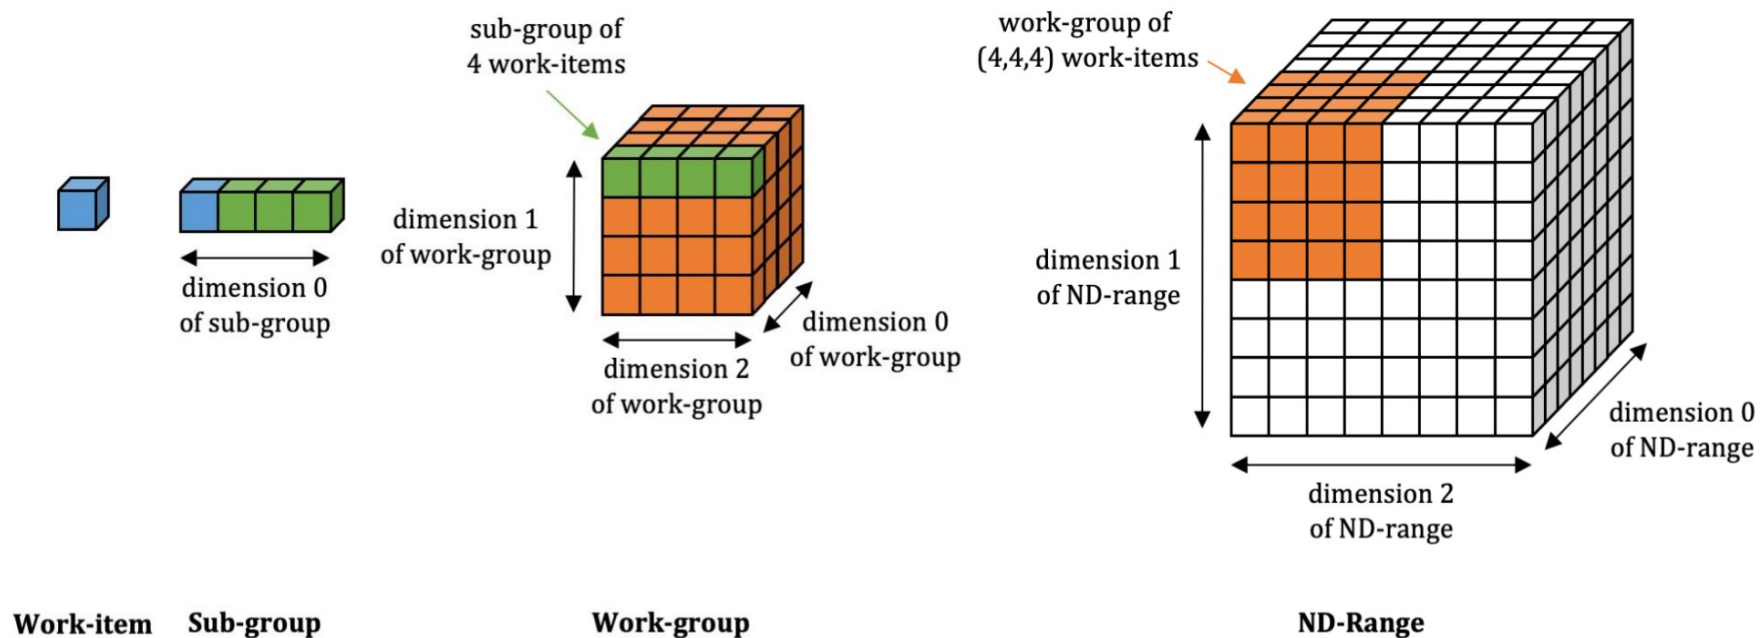
Write a SYCL program using Sub-Group and **group algorithms** to accomplish computation

## Sub-groups

Sub-groups are a **subset of the work-items** that are executed Simultaneously or with additional scheduling guarantees.

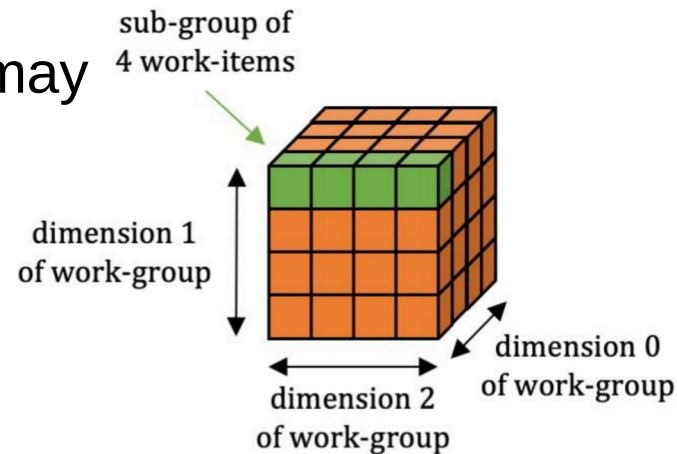
Leveraging sub-groups will help to **map execution to low level hardware** and may help in achieving **higher performance**.

# Sub-groups



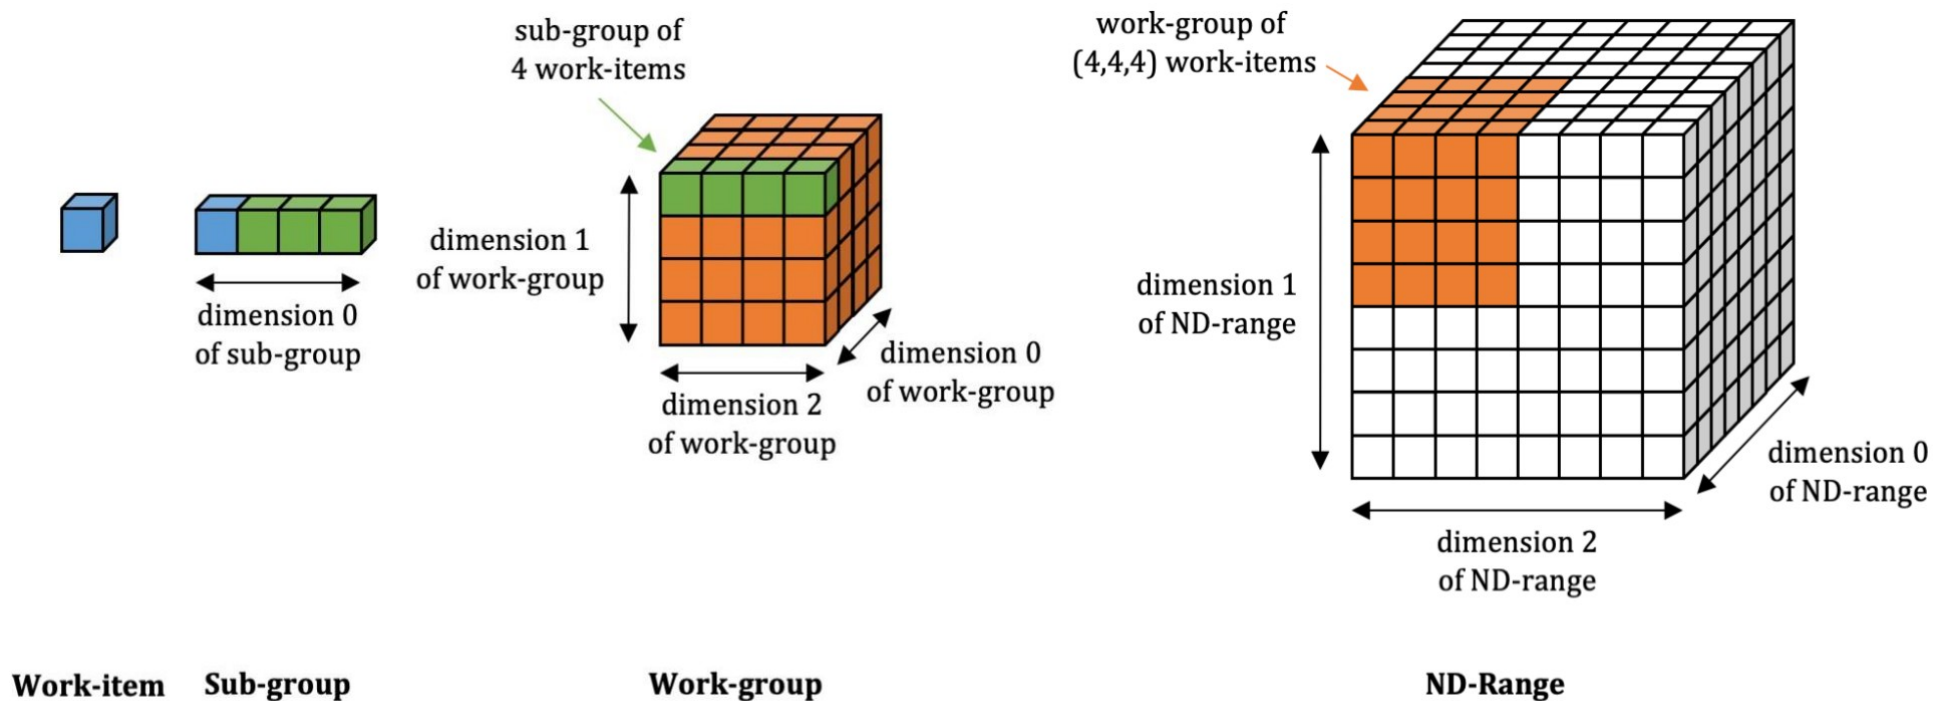
# Sub-groups

- A subset of work-items withing a work-group that may **map to vector hardware.**
- Why use sub-groups?
  - Work-items in a sub\_group can communicate directly using shuffle operations
  - Work-items in a sub\_group can synchronize using sub\_group barriers and guarantee memory consistency using sub\_group memory fences
  - Work-items in a sub\_group have access to sub\_group collectives, providing fast implementations of common parallel patterns.



# Sub-groups

- Sub-group = subset of work-items withing a work-group
- Parallel execution with `ND-RANGE` kernel helps to get access to work-group and sub-group



# Sub-groups

```
h.parallel_for(nd_range<1>(N,B), [=](nd_item<1> item)
{
    auto sg = item.get_sub_group();

    // KERNEL CODE
});
```

sub\_group class

- The sub-group handle can be obtained from the nd\_item using the `get_sub_group()`.
- Once you have the sub-group handle, you can **query** for more information about the sub-group, do **shuffle** operations or use **collective** functions.
- Explicit kernel attribute  
[[ intel::reqd\_sub\_group\_size(N) ]]  
to control the sub-group size

# Sub-groups

The sub-group handle can be queried to get other information:

- `get_local_id()` returns the index of the work-item within its sub-group
- `get_local_range()` returns the size of sub\_group
- `get_group_id()` returns the index of the sub-group
- `get_group_range()` returns the number of sub-groups within the parent work-group

```
h.parallel_for(nd_range<1>(N,B), [=](nd_item<1> item){
    auto sg = item.get_sub_group();

    if(sg.get_local_id() == 0){
        out << "sub_group id: " << sg.get_group_id()[0]
            << " of " << sg.get_group_range()
            << ", size=" << sg.get_local_range()[0]
                << endl;
    }
});
```

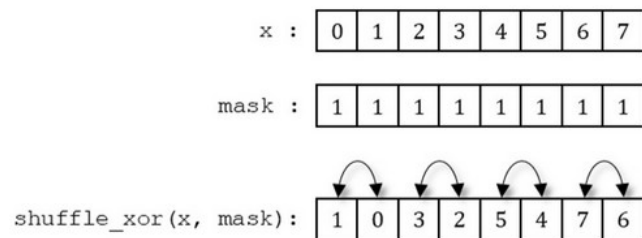
```
sub_group id: 1 of 4, size=16
sub_group id: 3 of 4, size=16
sub_group id: 2 of 4, size=16
sub_group id: 0 of 4, size=16
```



# Sub-group Shuffles

- One of the most useful features of sub-groups is the ability to communicate directly between individual work-items **without explicit memory operations**.
- Shuffle operations enable us to remove work-group **local memory usage** from our kernels and/or to avoid unnecessary repeated accesses to global memory.

```
h.parallel_for(nd_range<1>(N,B), [=](nd_item<1> item){  
    auto sg = item.get_sub_group();  
    size_t i = item.get_global_id(0);  
  
    /* Shuffles */  
    //data[i] = sg.shuffle(data[i], 2);  
    //data[i] = sg.shuffle_up(0, data[i], 1);  
    //data[i] = sg.shuffle_down(data[i], 0, 1);  
    data[i] = sg.shuffle_xor(data[i], 1);  
});
```



# Sub-group Collectives

- The collective functions provide implementations of closely- related **common parallel patterns**.
- Providing these implementations as library functions **increases developer productivity** and gives implementations the ability to generate highly optimized code for individual target devices.

```
h.parallel_for(nd_range<1>(N,B), [=](nd_item<1> item){
    auto sg = item.get_sub_group();

    size_t i = item.get_global_id(0);

    /* Collectives */
    data[i] = reduce(sg, data[i], plus<>());
    //data[i] = reduce(sg, data[i], std::maximum<>());
    //data[i] = reduce(sg, data[i], std::minimum<>());
});
```

# Sub Groups

## Sub-Group Group Algorithms

- Group algorithms provide implementations of closely-related **common parallel patterns**.
- Providing implementations as library functions **increases developer productivity** and gives implementations the ability to generate highly optimized code for individual target devices.

```
h.parallel_for(nd_range<1>(N,B), [=](nd_item<1> item)
{
    auto sg = item.get_sub_group();
    size_t i = item.get_global_id(0);

    /* Collectives */
    data[i] = reduce(sg, data[i], plus<>());
    //data[i] = reduce(sg, data[i], maximum<>());
    //data[i] = reduce(sg, data[i], minimum<>());
});
```

# Specifying the Sub-Group Size

The sub-group size can be **configured separately** for each kernel.

The set of available sub-group sizes is **hardware-specific**.

```
q.parallel_for(range<1>(N),  
               [=](id<1> id) [[intel::reqd_sub_group_size(16)]]  
{  
    // KERNEL CODE  
});
```

The sub-group size can be tuned even for kernels that do not use the **sub\_group** class (e.g. to tune for SIMD width and register usage).

# Sub Groups

- **Summary**
  - What are Sub-Groups?
  - Why are they useful?
  - Learned about sub-group shuffle operations and using sub-group collectives

# Reductions

A reduction produces a **single value by combining multiple values** in an unspecified order.

- **Parallelizing reductions** can be tricky because of the nature of computation and accelerator hardware.
- SYCL 2020 introduces a **simplified** approach for reductions in heterogenous programming

# Simple Reduction

Let's look a simple reduction example:  
***Addition of  $N$  items***

A simple **for-loop** in kernel function can accomplish reduction.

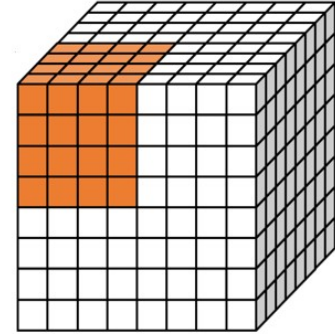
But, for-loop is **not efficient** and does not take advantage of parallelism in hardware.

```
queue q;  
int *data = malloc_shared<int>(N, q);  
for (int i = 0; i < N; i++) data[i] = i;  
  
q.single_task([=]()  
{  
    int sum = 0;  
    for(int i = 0; i < N; i++)  
    {  
        sum += data[i];  
    }  
    data[0] = sum;  
}).wait();  
  
std::cout << "Sum = " << data[0] << std::endl;
```

# Parallelizing Reductions



**work-group** executions are mapped to Compute Units on hardware.



Reduction can be parallelized by first reducing items in each work-group using ND-range kernel, **multiple work-groups can execute in parallel** depending on number of compute units on hardware.



# Work-Group Reduction

**ND-Range** kernel can be used to compute sum of all items in each work-group

**reduce()** function will simplify reduction of items in a work-group

A simple **for-loop** in `single_task` kernel function can then accomplish final reduction of each work-group sums.

```
q.parallel_for(nd_range<1>(N, B), [=](nd_item<1> item)
{
    auto wg = item.get_group();
    size_t i = item.get_global_id(0);

    /// Adds all elements in work_group using work_group reduce
    int sum_wg = reduce(wg, data[i], plus<>());

    /// write work_group sum to first location for each work_group
    if (item.get_local_id(0) == 0) data[i] = sum_wg;
});
```

```
q.single_task([=]()
{
    int sum = 0;
    for(int i=0; i<N; i+=B){
        sum += data[i];
    }
    data[0] = sum;
});
```

*Some parallelism  
achieved but  
code is still  
complex with 2  
kernel functions*

# Simplified Reduction

SYCL 2020 introduces  
reduction object in  
`parallel_for`

**reduction** object in  
`parallel_for` encapsulates the  
reduction variable, an  
optional operator identity  
and the reduction operator.

*Removes the need for two  
step approach using two  
kernel functions.*

```
queue q;  
auto data = malloc_shared<int>(N, q);  
for (int i = 0; i < N; i++) data[i] = i;  
  
auto sum = malloc_shared<int>(1, q);  
sum[0] = 0;  
  
q.parallel_for(nd_range<1>{N, B},  
               reduction(sum, plus<>()), [=](nd_item<1> it, auto& sum)  
{  
    int i = it.get_global_id(0);  
    sum += data[i];  
}).wait();  
  
std::cout << "Sum = " << sum[0] << std::endl;
```

# Multiple Reductions in one kernel

```
myQueue.submit([&](handler& cgh)
{
    // Input values to reductions are standard accessors (or USM pointers)
    auto inputValues = accessor(valuesBuf, cgh);

    // Create temporary objects describing variables with reduction semantics
    auto sumReduction = reduction(sumBuf, cgh, plus<>());
    auto maxReduction = reduction(maxBuf, cgh, maximum<>());

    // parallel_for performs two reduction operations
    cgh.parallel_for(range<1>{1024}, sumReduction, maxReduction,
        [=](id<1> idx, auto& sum, auto& max)
        {
            sum += inputValues[idx];
            max.combine(inputValues[idx]);
        });
});
```

# Useful Links

## Open source projects

oneAPI Data Parallel C++ compiler: [github.com/intel/llvm](https://github.com/intel/llvm)

Graphics Compute Runtime: Graphics [github.com/intel/compute-runtime](https://github.com/intel/compute-runtime)

Compiler: [github.com/intel/intel-graphics-compiler](https://github.com/intel/intel-graphics-compiler)

SYCL 2020: [tinyurl.com/sycl2020-spec](https://tinyurl.com/sycl2020-spec)

DPC++ Extensions: [tinyurl.com/dpcpp-ext](https://tinyurl.com/dpcpp-ext)

Environment Variables: [tinyurl.com/dpcpp-env-vars](https://tinyurl.com/dpcpp-env-vars)

DPC++ book: [tinyurl.com/dpcpp-book](https://tinyurl.com/dpcpp-book)

SYCL Academy [github.com/codeplaysoftware/syclacademy/tree/main](https://github.com/codeplaysoftware/syclacademy/tree/main)

Code samples:  
[github.com/intel/llvm/tree/sycl/sycl/test](https://github.com/intel/llvm/tree/sycl/sycl/test)  
[github.com/intel/llvm/tree/sycl/sycl/test-e2e](https://github.com/intel/llvm/tree/sycl/sycl/test-e2e)  
[github.com/oneapi-src/oneAPI-samples](https://github.com/oneapi-src/oneAPI-samples)

# Hands-on Exercises

## SYCL Lab 2 - Unified Shared Memory

# Notices and Disclaimers

- Performance varies by use, configuration and other factors. Learn more at [www.Intel.com/PerformanceIndex](http://www.Intel.com/PerformanceIndex)
- Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See backup for configuration details. No product or component can be absolutely secure.
- You may not use or facilitate the use of this document in connection with any infringement or other legal analysis concerning Intel products described herein. You agree to grant Intel a non-exclusive, royalty-free license to any patent claim thereafter drafted which includes subject matter disclosed herein.
- The products described may contain design defects or errors known as errata which may cause the product to deviate from published specifications. Current characterized errata are available on request.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document, with the sole exception that a) you may publish an unmodified copy and b) code included in this document is licensed subject to the Zero-Clause BSD open source license (0BSD), <https://opensource.org/licenses/0BSD>. You may create software implementations based on this document and in compliance with the foregoing that are intended to execute on the Intel product(s) referenced in this document. No rights are granted to create modifications or derivatives of this document.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document, with the sole exception that code included in this document is licensed subject to the Zero-Clause BSD open source license (0BSD), <http://opensource.org/licenses/0BSD>.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.
- Code names are used by Intel to identify products, technologies, or services that are in development and not publicly available. These are not "commercial" names and not intended to function as trademarks.
- © Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.

# Back Up

## Details about Intel® oneAPI Toolkits

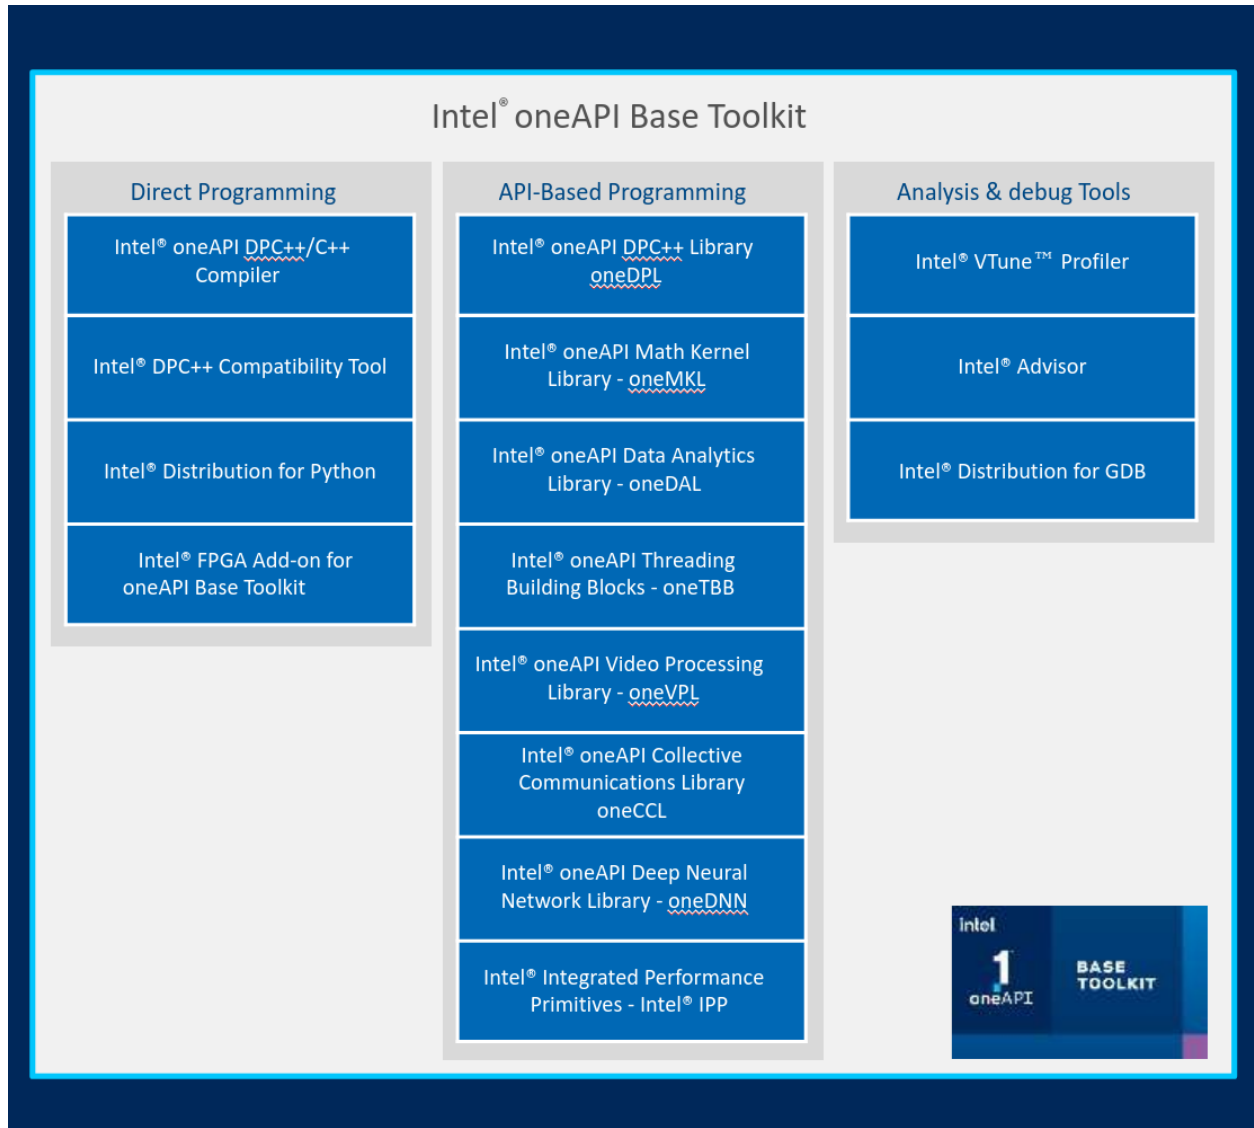
# Intel® oneAPI Base Toolkit

## Accelerate Data-centric Workloads

A core set of core tools and libraries for developing high-performance applications on Intel® CPUs, GPUs, and FPGAs.

### Who Uses It?

- A broad range of developers across industries
- Add-on toolkit users since this is the base for all toolkits





# Intel® oneAPI Base Toolkit

Accelerate Data-centric Workloads

## Top Features/Benefits

- Data Parallel C++ compiler, library and analysis tools
- SYCLomatic / DPC++ Compatibility tool helps migrate CUDA code to C++ with SYCL
- Python distribution includes accelerated scikit-learn, NumPy, SciPy libraries
- Optimized performance libraries for threading, math, data analytics, deep learning, and video/image/signal processing

[Learn More](#)

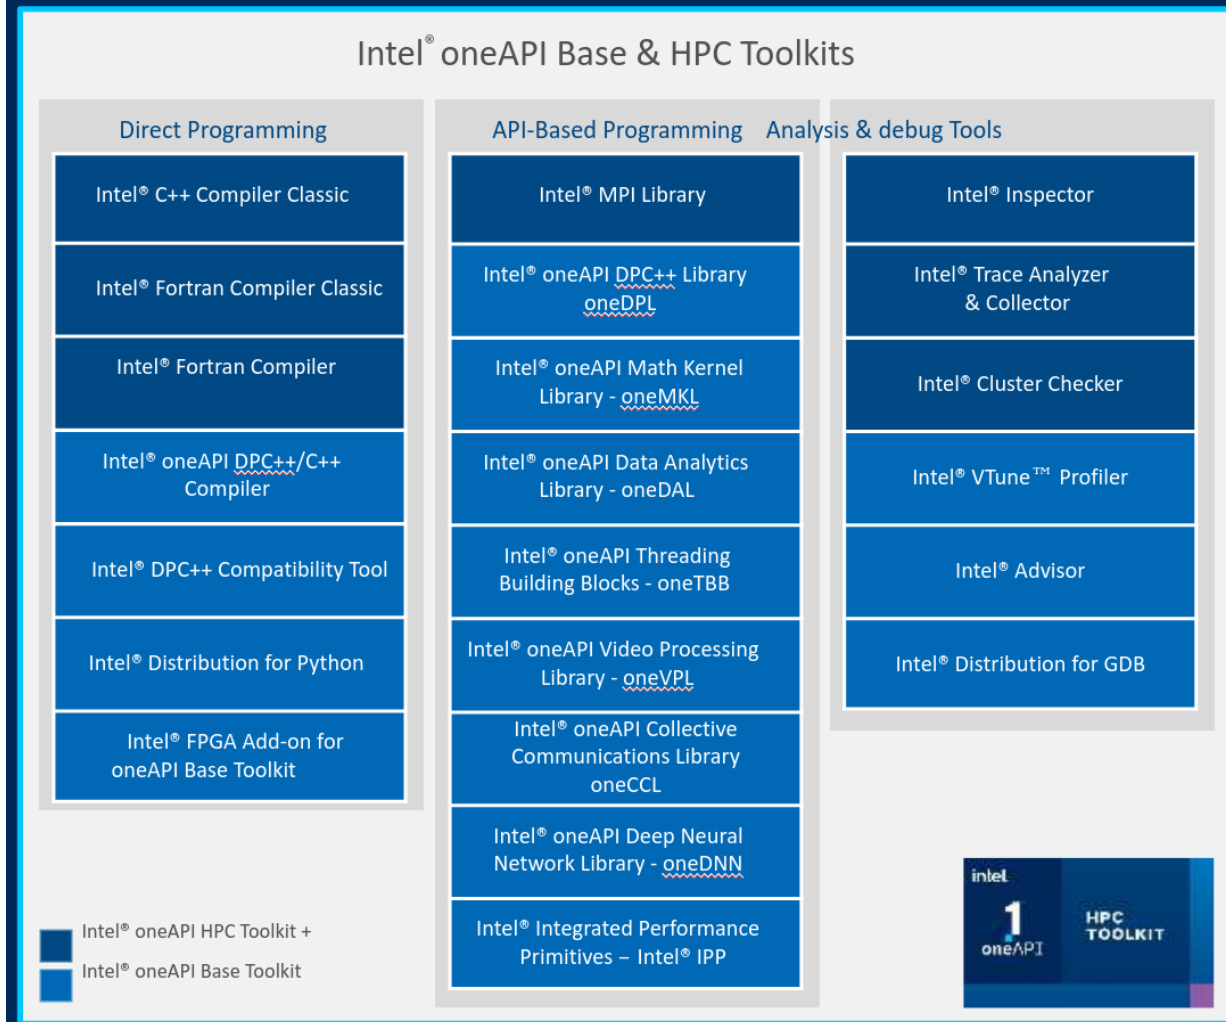
# Intel® oneAPI HPC Toolkit

## Accelerate Data-centric Workloads

A core set of core tools and libraries for developing high-performance applications on Intel® CPUs, GPUs, and FPGAs.

## Who Uses It?

- A broad range of developers across industries
- Add-on toolkit users since this is the base for all toolkits



# Intel® oneAPI HPC Toolkit

Accelerate Data-centric Workloads

## Top Features/Benefits

- Data Parallel C++ compiler, library and analysis tools
- SYCLomatic / DPC++ Compatibility tool helps migrate CUDA code to C++ with SYCL
- Python distribution includes accelerated scikit-learn, NumPy, SciPy libraries
- Optimized performance libraries for threading, math, data analytics, deep learning, and video/image/signal processing

# Summary

- oneAPI cross-architecture, one source programming model provides freedom of XPU choice.  
Apply your skills to the next innovation, not to rewriting software for the next hardware platform.
- Intel® oneAPI Toolkit products take full advantage of accelerated compute by maximizing performance across Intel CPUs, GPUs, and FPGAs.
- Develop confidently with a proven set of cross-architecture libraries and advanced tools that interoperate with existing performance programming models.



# Intel® oneAPI HPC Toolkit

Accelerate Data-centric Workloads

## Top Features/Benefits

- Data Parallel C++ compiler, library and analysis tools
- SYCLomatic / DPC++ Compatibility tool helps migrate CUDA code to C++ with SYCL
- Python distribution includes accelerated scikit-learn, NumPy, SciPy libraries
- Optimized performance libraries for threading, math, data analytics, deep learning, and video/image/signal processing